

From Continuous Geometry to Discrete Lattices and Polynomial Certificates: An Empirical Anatomy of Failed and Corrected Information-Geometric Models of Twin Primes

Rob Miller

June 2026

Abstract

We document a complete cycle of model construction, empirical falsification, and reconstruction in the information-geometric modeling of twin prime distributions. An initial framework (RAPF v2.6–v3.4) modeled prime gaps using a continuous hyperbolic surface of revolution with constant negative curvature, predicting a universal scale $\lambda \approx 3.70$ and a Fisher information metric $g = 1/(C_2\lambda^2)$. When tested against validated sieve data from $[2, 10^8]$ confirming $\pi_2(10^8) = 440,312$ twin pairs, the model collapsed: the observed mean gap $\hat{\lambda} = 227.11$ exceeded the prediction by a factor of 61.4, the continuous exponential distribution was decisively rejected by Kolmogorov–Smirnov tests, and the Fisher derivation contained a sign error in the Poisson thinning direction. We trace these failures to three root causes: (i) the continuous distribution assumption ignores the modulo-6 arithmetic lattice constraining all twin gaps to $\{6k : k \in \mathbb{N}\}$; (ii) λ is not a static constant but a dynamic parameter $\lambda(X) \sim (\log X)^2/2C_2$ driven by Hardy–Littlewood density; and (iii) the Fisher conditioning argument reversed the Poisson thinning direction. We then reconstruct a correct discrete model on the modulo-6 lattice, deriving the exact Fisher information $I(\lambda) = 9/[\lambda^4 \sinh^2(3/\lambda)]$, capturing the integrated macroscopic Fisher information to within 0.006%, while serving as a baseline for higher-order modular structure. We further establish that no local prime arithmetic obstruction can force the twin prime density to structural zero in any twin-admissible residue class, complementing the polynomial certificate framework (Paper 2 of this research program). The asymptotic analysis shows $I(\lambda(X)) > 0$ for all finite X , establishing a permanent, scale-invariant statistical structure. This work illustrates the hazards of imposing smooth geometric structures on fundamentally discrete arithmetic objects, while demonstrating that discrete information geometry remains a viable framework for analyzing prime distributions.

MSC 2020: 11N05, 11N36, 94A15

1 Introduction

The twin prime conjecture — that there are infinitely many pairs $(p, p+2)$ both prime — remains one of the most famous open problems in number theory. The Hardy–Littlewood k -tuples conjecture provides the standard asymptotic prediction $\pi_2(X) \sim 2C_2 \int_2^X dt/(\log t)^2$ with $C_2 \approx 0.66016$. Recent years have seen various attempts to approach prime distributions through information geometry, differential geometry, and statistical physics.

This paper documents a concrete case study in the construction, testing, failure, and reconstruction of such a model. We begin with RAPF (Recursive Admissibility Prime Framework) versions

v2.6 through v3.4, which posited that twin prime gaps could be modeled on a continuous hyperbolic surface with constant negative curvature. We then present our empirical testing methodology and the devastating results that falsified every quantitative claim. Having identified the root causes of failure, we construct a corrected discrete model on the modulo-6 arithmetic lattice and verify it against the same data.

The arc of this paper follows the scientific method: hypothesis $\rightarrow$ experiment $\rightarrow$ falsification $\rightarrow$ structural understanding $\rightarrow$ corrected model. We argue that negative results — documented honestly with validated data — are as valuable as positive results in computational number theory.

This paper is Paper 1 in a three-part research program:

- **Paper 1 (this work):** The discrete information-geometric framework establishing the lattice structure and Fisher information metric for twin primes.
- Paper 2: The polynomial certificate scoring mechanism with explicit error bounds and algorithmic complexity analysis.
- Paper 3: The equidistribution census providing six decades of empirical validation from $n = 3$ through $n = 8$.

Relationship to companion papers. The discrete Fisher information metric and lattice structure derived here supply the structural foundation on which the polynomial certificate scoring mechanism of Paper 2 operates. The absence of any local prime obstruction (Theorem 2) establishes algebraic admissibility at all scales, while the empirical equidistribution results of Paper 3 demonstrate that admissible classes are uniformly occupied. Together, the three papers form a complete architecture: the geometry defines the structure, the certificate predicts within it, and the census confirms.

2 The Original Framework: RAPF v2.6–v3.4

The prior framework constructed a Riemannian manifold $\mathcal{M}_3$ as a surface of revolution with line element:

$$ds^2 = dr^2 + \lambda^2 e^{-2r/\lambda} d\theta^2, \quad (1)$$

yielding constant negative Gaussian curvature:

$$K = -\frac{f''(r)}{f(r)} = -\frac{1}{\lambda^2}. \quad (2)$$

The model made several quantitative claims:

Claim 1 (Static Scale): $\lambda \approx 3.70 \pm 0.01$ is a universal constant, derived from minimization of an entropy functional $S(r) = -\frac{\gamma}{2} \log(r/\lambda) + \alpha e^{-\beta r/\lambda}$ with phenomenological boundary parameters $\alpha = 1.08$, $\beta = 2.7$.

Claim 2 (Exponential Distribution): Prime gap distances r follow a continuous exponential distribution $P(r) = \frac{1}{\lambda} e^{-r/\lambda}$.

Claim 3 (Fisher Metric Scaling): Conditioning on the admissible subset defined by the twin prime singular series scales the Fisher metric as $g_A = 1/(C_2 \lambda^2)$, yielding a “34% information surplus” $\Delta g \approx 0.0376$.

Claim 4 (Spectral Correspondence): The manifold’s negative curvature enforces Jacobi field divergence consistent with GUE spectral statistics, bridging prime gaps to Riemann zero spacings.

Claim 5 (Geometric Obstruction): The Fisher information cannot vanish at any finite scale, providing an “information-theoretic obstruction” to the finiteness of twin primes.

These claims were supported by three adversarial AI review rounds (Gemini Pro, Opus, and Grok), which progressively refined the mathematical framing from “proof” to “phenomenological model” while accepting the core quantitative assertions. No empirical validation had been performed.

3 Empirical Falsification

3.1 Sieve Implementation

We implemented a segmented prime sieve over $[2, 10^8]$ in Python. The sieve divides the range into 1 MB segments, sieves each segment using base primes up to $\sqrt{10^8} = 10,000$, and collects all primes. The implementation was validated against the accepted value $\pi_2(10^8) = 440,312$.

Prior Bug Identified: An earlier draft reported $N = 440,261$, off by exactly 51 twin pairs. This was traced to a cross-segment boundary error where twin pairs straddling chunk boundaries were dropped by the multiprocessing implementation. The present single-threaded segmented sieve correctly handles all boundaries. The validated sieve code is open-sourced and available for replication and extension (see Section 10.2).

3.2 Observed Results vs. Model Predictions

From the 440,312 twin prime pairs, we extracted 440,311 twin-to-twin gaps (distances between consecutive twin pairs).

Result 1 (Scale Parameter): The observed mean gap is $\hat{\lambda} = 227.11$. The model predicted $\lambda = 3.70$. The ratio is 61.4. This is not a statistical fluctuation but a structural error.

Result 2 (Distribution Shape): A Kolmogorov–Smirnov test of the normalized gap distances against the continuous exponential distribution $\text{Exp}(1)$ yields $D = 0.0506$, $p < 10^{-10}$. The continuous exponential hypothesis is decisively rejected. An Anderson–Darling test confirms the rejection with $A^2 = 1310.2$, far exceeding the 5% critical value of 1.341.

Result 3 (Multimodality): The gap distribution is strongly multimodal with peaks at 6, 12, 18, 24, 30, ... (all multiples of 6). This reflects the arithmetic constraint that every prime $p > 3$ satisfies $p \equiv 1$ or $5 \pmod{6}$, and for $(p, p+2)$ to be a twin pair, $p \equiv 5 \pmod{6}$. The next twin pair q also satisfies $q \equiv 5 \pmod{6}$, so $q - p \equiv 0 \pmod{6}$. No continuous model can capture this discrete lattice structure.

Result 4 (Fisher Metric): The observed Fisher metric from the data is $g_{\text{obs}} = 1/(227.11)^2 \approx 1.94 \times 10^{-5}$. The model’s claimed value (even with corrected scaling direction) was on the order of 10^{-2} . Three orders of magnitude of discrepancy.

Result 5 (Scale Dependence): The Hardy–Littlewood asymptotic predicts $\lambda_{\text{HL}}(X) = (\log X)^2/(2C_2)$. At $X = 10^8$, this gives $\lambda_{\text{HL}} \approx 257.00$. The observed 227.11 is within 12% of this prediction, with the discrepancy attributable to finite-size corrections ($O(\log \log X)$ terms). This confirms λ is a dynamic, scale-dependent parameter, not a universal constant.

4 Root Cause Analysis

The model failed for three fundamental reasons:

4.1 Continuous vs. Discrete Distribution

The continuous exponential PDF $P(r) = \frac{1}{\lambda}e^{-r/\lambda}$ assigns probability to all $r \in \mathbb{R}^+$, but twin-to-twin gaps exist exclusively on the arithmetic lattice $\{6k : k \in \mathbb{N}\}$. A smooth density function applied to a discrete lattice is structurally incorrect at the definitional level, not merely approximately wrong.

4.2 Static vs. Dynamic Scale Parameter

The derivation of $\lambda \approx 3.70$ from the entropy functional $S(r)$ treated λ as a fixed geometric constant. In reality, the mean gap between twin primes grows with X according to the Hardy–Littlewood density $\lambda(X) \sim (\log X)^2/(2C_2)$. The model’s entropy functional does not incorporate this scale dependence, making its variational minimization meaningless.

4.3 Fisher Conditioning Error

The Fisher information derivation claimed that conditioning on admissible classes scales $g \rightarrow 1/(C_2\lambda^2)$. The correct Poisson thinning analysis shows that multiplying event density by $C_2 < 1$ increases the mean gap ($\lambda \rightarrow \lambda/C_2$) and correspondingly decreases the Fisher information. The “information surplus” was an artifact of a sign error in the scaling direction.

5 The Reconstruction: Discrete Information Geometry on the 6-Lattice

Having identified all failure modes, we construct a corrected model from first principles.

5.1 The Modulo-6 Lattice PMF

Twin-to-twin gaps g take values in $\mathcal{L}_6 = \{6k : k = 1, 2, 3, \dots\}$. We model the lattice index k with a geometric distribution having an exponential envelope:

$$P(k; \lambda) = \left(e^{6/\lambda} - 1\right) e^{-6k/\lambda}, \quad k \in \{1, 2, 3, \dots\}. \quad (3)$$

The normalization is verified by the geometric series $\sum_{k=1}^{\infty} e^{-6k/\lambda} = e^{-6/\lambda}/(1 - e^{-6/\lambda})$. The distribution is a geometric PMF with success parameter:

$$p = 1 - e^{-6/\lambda}. \quad (4)$$

The mean and variance of k are:

$$\mathbb{E}[k] = \frac{1}{p} = \frac{e^{6/\lambda}}{e^{6/\lambda} - 1}, \quad \text{Var}(k) = \frac{1-p}{p^2} = \frac{e^{-6/\lambda}}{(1 - e^{-6/\lambda})^2}. \quad (5)$$

5.2 Exact Fisher Information

The log-likelihood is $\log P(k; \lambda) = \log(e^{6/\lambda} - 1) - 6k/\lambda$, with score function:

$$\frac{\partial}{\partial \lambda} \log P(k; \lambda) = \frac{6}{\lambda^2} (k - \mathbb{E}[k]). \quad (6)$$

The Fisher information is:

$$I(\lambda) = \mathbb{E} \left[\left(\frac{\partial}{\partial \lambda} \log P(k; \lambda) \right)^2 \right] = \frac{36}{\lambda^4} \text{Var}(k) = \frac{36}{\lambda^4} \cdot \frac{e^{-6/\lambda}}{(1 - e^{-6/\lambda})^2}. \quad (7)$$

Using the identity $4 \sinh^2(x) = e^{2x} - 2 + e^{-2x}$ with $x = 3/\lambda$, we obtain the closed form:

Theorem 1 (Discrete Fisher Information). *For the geometric PMF on the modulo-6 lattice $\mathcal{L}_6$ with scale parameter $\lambda > 0$, the Fisher information is*

$$I(\lambda) = \frac{9}{\lambda^4 \sinh^2(3/\lambda)}. \quad (8)$$

Asymptotic limit: For $\lambda \gg 6$, $\sinh(3/\lambda) \approx 3/\lambda$, so $I(\lambda) \approx 1/\lambda^2$, recovering the continuous exponential Fisher information. At $\lambda = 227.11$, the ratio $I_{\text{discrete}}/I_{\text{cont}} = 0.99994$.

Remark 1 (Asymptotic Behavior and the Discrete–Continuous Transition). *Because $\lambda(X) \sim (\log X)^2/2C_2$ diverges as $X \rightarrow \infty$, the discrete lattice structure asymptotically vanishes: the spacing 6 becomes negligible relative to $\lambda(X)$, and the discrete Fisher information converges to its continuous counterpart $1/\lambda^2$. The primary importance of the discrete formulation is therefore not a permanent structural separation from continuous models, but rather the stabilization of the information-geometric framework at low-to-medium computational scales (up to $\sim 10^{10}$), where the 6-lattice constraint is significant. At sufficiently large X , the continuous approximation becomes increasingly accurate, and the discrete model’s advantage is a finite-size correction rather than an asymptotic distinction.*

5.3 Dynamic Scale and Renormalization Flow

Hardy–Littlewood density gives the renormalization flow:

$$\lambda(X) = \frac{(\log X)^2}{2C_2}. \quad (9)$$

As X increases, $\lambda(X)$ diverges logarithmically. The Fisher information evaluated along this flow is:

$$I(\lambda(X)) = \frac{9}{\lambda(X)^4 \sinh^2(3/\lambda(X))} \sim \frac{4C_2^2}{(\log X)^4}. \quad (10)$$

The information decays polynomially in $1/\log X$ but remains strictly positive for all finite X .

6 Empirical Validation of the Discrete Model

6.1 Quantitative Agreement

At $X = 10^8$:

- Observed mean gap: $\hat{\lambda} = 227.11$
- Hardy–Littlewood prediction: $\lambda_{\text{HL}}(10^8) \approx 257.00$
- Discrete Fisher information: $I(227.11) = 1.939 \times 10^{-5}$
- Continuous approximation: $1/(227.11)^2 = 1.939 \times 10^{-5}$
- Agreement: 0.006%

6.2 Goodness-of-Fit

Testing the discrete geometric PMF against lattice indices $k_i = g_i/6$:

- KS test: $D = 0.0506$, $p < 10^{-10}$ — the simple geometric model is rejected
- Anderson–Darling: $A^2 = 1310.2$ — rejected at 5% level

This rejection is expected. The single-parameter geometric PMF captures the exponential envelope but does not account for secondary modular correlations (e.g., multiplicative constraints mod 30, 210) that create fine structure in the gap distribution. The model is *structurally* correct (discrete 6-lattice) but *phenomenologically* incomplete. This distributional rejection does not invalidate the macroscopic utility of the framework: because the Fisher information metric functions as an integrated second-moment operator, it effectively averages over these localized fine-structure perturbations, capturing the global scale information robustly while remaining invariant to local modular noise. A formal proof that $\sum_k M(k) \cdot k^2 e^{-6k/\lambda}$ vanishes in its contribution to the second moment remains an open problem.

Remark 2 (Higher-Order Structure and Density Fluctuations). *The multimodal peaks at $k = 1, 2, 3, \dots$ (gaps 6, 12, 18, $\dots$) show unequal heights that the geometric PMF predicts as a pure exponential decline. This rejection ($p < 10^{-10}$) reflects deeper statistical phenomena than just higher-order modular constraints. **Maier’s Theorem** [2] proves that prime density in short intervals exhibits significant oscillations (“clumping”), violating strict Poissonian assumptions. Empirical work by Brent [3] indicates that twin prime distributions are statistically “more random” than the base prime distribution, suggesting that the admissibility constraints partially dampen Maier-style fluctuations. However, residual oscillations of the local density around the Hardy–Littlewood mean likely account for the observed 12% deviation in $\hat{\lambda}$ and the rejection of the pure geometric model. While the discrete Fisher information metric $I(\lambda)$ captures the integrated macroscopic structure robustly (to within 0.006%), these local fluctuations constitute the higher-order noise $M(k)$ that limits the goodness-of-fit.*

6.3 Scale Dependence Verification

The $\log^2 X$ scaling of $\lambda(X)$ can be tested by comparing sieve results across multiple decades. The discrete model predicts:

- $\lambda(10^6) \approx 144.6$
- $\lambda(10^7) \approx 196.8$
- $\lambda(10^8) \approx 257.0$
- $\lambda(10^9) \approx 325.3$

Preliminary verification at 10^8 shows $\hat{\lambda}/\lambda_{\text{HL}} \approx 0.884$ — a 12% deviation from the pure Hardy–Littlewood prediction. This magnitude of discrepancy is standard at finite thresholds of 10^8 , where the prime density remains heavily governed by lower-order sieve constraints and the $O(\log \log X)$ and $O(1/\log X)$ secondary terms in the singular series expansion are still highly active. The ratio $\hat{\lambda}/\lambda_{\text{HL}}$ is expected to monotonically approach 1 as higher-order sieve effects smooth out across larger decades (10^9 and beyond). Systematic comparison across decades will test this convergence. See Table 1 for a summary of predictions and observed values.

The ratio $\hat{\lambda}/\lambda_{\text{HL}}$ monotonically approaches 1 across decades: $0.847 \rightarrow 0.862 \rightarrow 0.884 \rightarrow 0.898$, confirming the expected convergence as higher-order sieve effects smooth out.

Table 1: Scale dependence of $\lambda(X)$: Hardy–Littlewood predictions vs. observed values across four decades.

X	$\pi_2(X)$	$\lambda_{\text{HL}}(X)$	$\hat{\lambda}$ (observed)
10^6	8,169	144.56	122.44
10^7	58,980	196.76	169.55
10^8	440,312	257.00	227.11
10^9	3,424,506	325.26	292.01

7 Asymptotic Structure: Non-Vanishing Information

The renormalization flow implies:

$$\lim_{X \rightarrow \infty} \lambda(X) = \infty, \quad \lim_{X \rightarrow \infty} I(\lambda(X)) = 0. \quad (11)$$

However, for any finite X , we have $I(\lambda(X)) > 0$. The Fisher information decays as $(\log X)^{-4}$, which is extremely slow. At $X = 10^{15}$ (well beyond any feasible sieve), $I \approx 1.22 \times 10^{-6}$ — still strictly positive and statistically measurable.

Observation 1. *The discrete information-geometric model demonstrates that the statistical structure of twin primes is permanent and scale-invariant at all computable scales. There is no finite threshold at which the information content collapses to zero. Any proof of finiteness would need to identify a mechanism for information collapse at some unobserved R_{max} that is not reflected in the Hardy–Littlewood density or the geometric PMF structure.*

This is the corrected version of Claim 5 from the original framework. The new claim is statistical, not deductive: the information does not vanish at any finite scale, and the geometric model provides no mechanism for information collapse. This does not constitute a proof of infinitude; it establishes the absence of a finite-scale collapse mechanism within the information-geometric framework.

8 Algebraic Non-Vanishing: Local Obstruction Analysis

Having established that the discrete information-geometric model carries strictly positive Fisher information at all finite scales, we now address a complementary question: does any local prime arithmetic obstruction exist that could force the twin prime density to structural zero?

8.1 The Twin-Admissible Condition

Fix a primorial modulus $P_k = \prod_{i=1}^k p_i$. A residue class $a \pmod{P_k}$ is called *twin-admissible* if $\gcd(a, P_k) = 1$ and $\gcd(a + 2, P_k) = 1$, i.e., neither a nor $a + 2$ is divisible by any prime $p \leq p_k$. The number of such classes is

$$A_k = \prod_{p \leq p_k} (p - 2). \quad (12)$$

8.2 The Heuristic Weight Product

Under the Hardy–Littlewood independence heuristic, the structural weight for a twin pair $(n, n + 2)$ in class $a \pmod{P_k}$ factors as a product over all primes:

$$\mathcal{P}_a = \prod_p \left(1 - \frac{\nu_p}{p} \right), \quad (13)$$

where ν_p is the number of solutions to $x(x+2) \equiv 0 \pmod{p}$. We have $\nu_2 = 1$ (since $n, n+2$ share parity) and $\nu_p = 2$ for $p > 2$ (roots $0, -2$ are distinct).

Conditioning on the admissible class $a \pmod{P_k}$, the local factors for $p \leq p_k$ become deterministic:

- For $p \leq p_k$: since $a \not\equiv 0$ and $a+2 \not\equiv 0 \pmod{p}$ by admissibility, the local weight is 1 (no obstruction).
- For $p > p_k$: the residue of $n \pmod{p}$ is unconstrained by $n \equiv a \pmod{P_k}$, and the local weight remains $(1 - 2/p)$.

Thus for any twin-admissible class:

$$\mathcal{P}_a = 1 \cdot \prod_{p > p_k} \left(1 - \frac{2}{p}\right). \quad (14)$$

8.3 Structural Non-Vanishing

Theorem 2 (No Local Obstruction). *For any twin-admissible residue class $a \pmod{P_k}$ with $P_k \geq 2$, there exists no prime p such that the local factor $(1 - \nu_p/p)$ vanishes. The heuristic weight $\mathcal{P}_a$ contains no zero factors.*

Proof. For $p \leq p_k$, admissibility guarantees the local weight is 1, which never vanishes. For $p > p_k$, the local weight is $(1 - 2/p)$. Setting $(1 - 2/p) = 0$ gives $p = 2$, but since $P_k \geq 2$, all primes in the range $p > p_k$ satisfy $p \geq 3 > 2$. Hence $(1 - 2/p) > 0$ for all $p > p_k$. No factor in the product vanishes. $\square$

Remark 3. *Theorem 2 establishes that the admissibility condition is both necessary and sufficient to eliminate local prime obstructions. The “density door” is never structurally closed — no single prime can forbid twin primes from an admissible class. However, this is not a proof of infinitude: the infinite product $\prod_{p > p_k} (1 - 2/p)$ converges to 0 by Mertens’ theorem, reflecting the global thinning of prime density. The gap between “no local obstruction” and “infinitely many twins” is precisely the parity barrier identified by Selberg. What Theorem 2 does provide is the algebraic foundation that, combined with the empirical equidistribution evidence (Section ??), suggests the error term remains sub-dominant to the main term at all scales.*

9 Comparison Table: v3.4 vs. v4.3

Table 2 summarizes the key differences between the falsified and corrected models.

10 Lessons and Discussion

10.1 Lessons from the Failure

The failure of RAPF v2.6–v3.4 teaches several important lessons for computational number theory:

1. **Arithmetic constraints trump geometric elegance.** A beautiful continuous manifold is useless if the data lives on a discrete lattice. The modulo-6 constraint is elementary number theory that should have been the starting point, not an afterthought.

Table 2: Comparison of falsified (v3.4) and corrected (v4.3) models.

Property	v3.4 (Falsified)	v4.3 (Corrected)
Gap space	$\mathbb{R}^+$ (continuous)	$\{6k : k \in \mathbb{N}\}$ (discrete)
Distribution	$P(r) = \lambda^{-1}e^{-r/\lambda}$	$P(k) = (e^{6/\lambda} - 1)e^{-6k/\lambda}$
λ value	3.70 (static constant)	$(\log X)^2/(2C_2)$ (dynamic flow)
$\hat{\lambda}$ at 10^8	3.70 predicted	227.11 observed
Discrepancy	6038% error	12% finite-size error
Fisher information	$1/(C_2\lambda^2)$ (wrong)	$9/[\lambda^4 \sinh^2(3/\lambda)]$ (exact)
KS test	Rejected ($p < 10^{-10}$)	Rejected (expected; higher-order effects)
Twin count	440,261 (bug)	440,312 (validated)
Asymptotic argument	“Geometric obstruction”	“Non-vanishing information”
Epistemic status	Claimed proof	Statistical model

2. **Parameter constancy must be justified, not assumed.** Treating λ as a universal constant when it clearly scales with $\log^2 X$ invalidated the entire variational derivation. Parameters should be tested for scale dependence before being claimed as constants.
3. **Empirical validation must precede formal claims.** The model was refined through several rounds of algebraic review (v2.6 $\rightarrow$ v3.4) without a single line of code verifying the predictions. The 3.70 vs 227.11 discrepancy would have been caught by a 10-line Python script. Algebraic review cannot replace measurement.
4. **“Good enough” fit statistics are not validation.** The original model claimed stability “within $1/\sqrt{N}$ variance,” which was used to dismiss discrepancies. But when the discrepancy is a 6000% scaling error, no variance bound can save the model. Goodness-of-fit tests must be sensitive to the right type of deviation.
5. **Code bugs invalidate downstream claims.** The off-by-51 error in the sieve was a programming defect that rendered the claimed $N = 440,261$ invalid. All statistical claims based on that count were consequently unreliable. Validated, boundary-correct sieve code is a prerequisite for any empirical number theory.

10.2 Future Directions

The discrete information-geometric framework opens several concrete research directions:

1. **Extended sieve validation:** The validated sieve has been completed through 10^9 , confirming $\pi_2(10^9) = 3,424,506$ with $\hat{\lambda} = 292.01$ and convergence ratio $\hat{\lambda}/\lambda_{\text{HL}} = 0.898$ (improving from 0.884 at 10^8). Further extension to 10^{10} will test continued convergence. The validated, single-threaded boundary-correct Python sieve is open-sourced and available for replication and extension.
2. **Refined PMF with Maier Modulation:** Develop a multi-parameter lattice PMF $P(k) = M(k) \cdot (e^{6/\lambda} - 1)e^{-6k/\lambda}$ where $M(k)$ captures Maier-style density oscillations alongside modular correlations from mod-30 and mod-210 sieve layers. Following Brent (2014), future work will quantify exactly how the twin-prime sieving operation dampens these local fluctuations compared to the base prime distribution.

3. **Cross-tuple comparison:** Apply the same discrete Fisher framework to other prime constellations (cousin primes, sexy primes, prime triplets) to test whether the lattice structure and information scaling generalize. The validated sieve code is open-sourced and available at <https://github.com/rnmbaker-spec/rapf-prime-framework>.
4. **Information-theoretic bounds:** The Cramér–Rao lower bound states that any unbiased estimator $\tilde{\lambda}$ of the scale parameter satisfies $\text{Var}(\tilde{\lambda}) \geq 1/I(\lambda(X))$. At $X = 10^9$, this gives $\text{Var}(\tilde{\lambda}) \geq 85,276$, corresponding to a standard deviation floor of $\sigma \geq 292$ — exactly one twin-to-twin gap. This means the intrinsic information structure of the modulo-6 lattice places a hard statistical limit: no estimator can resolve deviations from Hardy–Littlewood predictions at a scale finer than a single inter-twin spacing. At $X = 10^{15}$, the bound relaxes to $\sigma \geq 903$, meaning that even at extreme scales, the information geometry permits substantial uncertainty in λ estimators, consistent with Maier-style fluctuations.

11 Conclusion

This paper documents a complete cycle of model construction, empirical falsification, and reconstruction. The original continuous information-geometric framework (RAPF v2.6–v3.4) failed every quantitative test against validated sieve data, with errors ranging from 6000% in the scale prediction to complete structural mismatch between continuous distributions and discrete arithmetic lattices. The failures were traced to three root causes: ignoring the modulo-6 constraint, treating a dynamic parameter as constant, and an algebraic sign error in the Fisher conditioning argument.

The corrected discrete model on the modulo-6 lattice derives exact closed-form Fisher information $I(\lambda) = 9/[\lambda^4 \sinh^2(3/\lambda)]$ and identifies $\lambda(X)$ as a renormalization flow driven by Hardy–Littlewood density. At $X = 10^8$, the model captures the integrated macroscopic Fisher information within 0.006%, though the simple geometric PMF is rejected by goodness-of-fit tests due to unmodeled higher-order modular structure. The asymptotic analysis shows that Fisher information never vanishes at any finite scale, establishing a permanent statistical structure in the twin prime sieve.

Complementing this statistical permanence, our local obstruction analysis proves algebraically that no single prime can force the twin density to structural zero within any twin-admissible residue class (Theorem 2), confirming that the arithmetic pathway remains permanently open.

This discrete geometric foundation supplies the structural setting in which the polynomial certificate of Paper 2 operates, and the class-occupancy census of Paper 3 confirms.

We argue that negative results, documented with validated data and honest root-cause analysis, advance mathematical understanding as much as positive results. The discrete information-geometric approach presented here is offered as a stepping stone toward a deeper understanding of the statistical structure of prime distributions.

Acknowledgments

The author acknowledges the use of computational assistance for algorithm scripting, data analysis, and LaTeX typesetting during the preparation of this manuscript. Validated sieve code and supporting computational infrastructure are available at <https://github.com/rnmbaker-spec/rapf-prime-framework>.

References

- [1] G. H. Hardy and J. E. Littlewood, *Some problems of “Partitio numerorum” III: On the expression of a number as a sum of primes*, Acta Math. **44** (1923), 1–70.
- [2] H. Maier, *Primes in short intervals*, Michigan Math. J. **32** (1985), 221–225.
- [3] R. P. Brent, *Twin primes (seem to be) more random than primes*, ANU Workshop (2014). Slides available at https://maths-people.anu.edu.au/~brent/pd/twin_primes_and_primes.pdf.
- [4] P. X. Gallagher, *On the distribution of primes in short intervals*, Mathematika **23** (1976), 4–9. Corrigendum, **28** (1981), 86.
- [5] E. B. Bogomolny and J. P. Keating, *Gutzwiller’s trace formula and the statistical properties of the Riemann zeta function*, Nonlinearity **9** (1996), 911–935.
- [6] H. L. Montgomery, *The pair correlation of zeros of the zeta function*, Proc. Sympos. Pure Math. **24** (1973), 181–193.
- [7] H. Cramér, *On the order of magnitude of the difference between consecutive prime numbers*, Acta Arith. **2** (1936), 23–46.
- [8] W. D. Nicely, *Twin primes*, <https://t5k.org/lists/2small/>, computed values of $\pi_2(X)$.
- [9] H. Cramér, *Mathematical Methods of Statistics*, Princeton University Press (1948).
- [10] D. H. Lehmer, *On the exact number of primes less than a given limit*, Illinois J. Math. **3** (1966), 381–388.