

Toward Equidistribution in Primorial Arithmetic Progressions: Numerical Results and Conjectures

Rob Miller

June 2026

Abstract

We investigate the distribution of primes in arithmetic progressions to primorial moduli $P_n = \prod_{i=1}^n p_i$. While the classical Bombieri–Vinogradov theorem gives no useful information for this sparse sequence of moduli, we observe that primorials are the canonical *smooth numbers* and connect this to a different analytical toolkit – subconvexity bounds and character sum estimates for smooth moduli (Irving 2016, Goldmakher 2010, Weyl differencing).

We complement this theoretical analysis with detailed computational evidence: for each primorial P_n with $n \leq 6$, we compute $\pi(P_n^2; P_n, a)$ for *every* twin-admissible residue class, extend with a 500-class random sample at $n = 7$, and perform the full census of all 378,675 twin-admissible classes at $n = 8$ ($P_8^2 \approx 9.4 \times 10^{13}$). The bias converges monotonically from 1.78% at $n = 3$ to below 0.001% at $n = 8$ — the structural density prediction from the RAPF framework (Paper 1) matches the actual integer count to within 0.0006% at $n = 8$. The Coefficient of Variation (CV) at $n = 8$ is 0.15%, suppressed $10\times$ below the Poisson random baseline, with zero empty classes observed among all 378,675 admissible residue classes.

We formulate three conjectures regarding primorial equidistribution and joint class occupancy. If proven, these would establish that primes fill admissible residue classes uniformly, with direct implications for the infinitude of twin primes.

1 Introduction

The Distribution of primes in arithmetic progressions

$$\pi(x; q, a) = \#\{p \leq x : p \equiv a \pmod{q}\} \quad (1)$$

is one of the central problems in analytic number theory. The Prime Number Theorem for arithmetic progressions asserts that for fixed q and $(a, q) = 1$,

$$\pi(x; q, a) \sim \frac{\text{li}(x)}{\phi(q)}. \quad (2)$$

The Bombieri–Vinogradov theorem [1, 2] extends this to an average over all moduli $q \leq Q$ with $Q \leq x^{1/2}(\log x)^{-B}$. However, the theorem says nothing about any *individual* modulus, and the sequence of primorials $P_n = \prod_{i=1}^n p_i$ is far too sparse (only $O(\log \log x)$ primorials up to x) for any sparse-set extension of Bombieri–Vinogradov [3, 4] to apply.

In this paper, we take a different approach. We observe that primorials are the canonical *smooth numbers*: every prime factor of P_n satisfies $p \leq p_n \sim \log P_n$. For smooth moduli, there exist published unconditional theorems that give *strictly better* bounds for Dirichlet L -functions and character sums than what holds for arbitrary moduli:

This paper is Paper 3 in a three-part research program:

- Paper 1: The discrete information-geometric framework establishing the lattice structure and Fisher information metric for twin primes.
- Paper 2: The polynomial certificate scoring mechanism with explicit error bounds and algorithmic complexity analysis.
- **Paper 3 (this work):** The equidistribution census providing six decades of empirical validation from $n = 3$ through $n = 8$.

Relationship to companion papers. The lattice structure and Fisher information metric established in Paper 1 define the arithmetic setting in which the polynomial certificate of Paper 2 operates. This paper provides the empirical census that validates both: demonstrating that primes distribute uniformly across the admissible classes defined in Paper 1, and supplying the benchmark datasets against which Paper 2’s predictions are tested. Together, the three papers form a complete architecture: the geometry defines the structure, the certificate predicts within it, and the census confirms.

- **Irving (2016)** [5]: $L(1/2, \chi) \ll q^{27/164+O(\delta)+\varepsilon}$ for smooth q , improving on the convexity bound $q^{1/4}$ and even the best general bound $q^{1/6+\varepsilon}$ of Petrow–Young [8].
- **Goldmakher (2010)** [6]: Character sums to smooth moduli satisfy a strengthened Pólya–Vinogradov inequality with smaller logarithmic factors.
- **Weyl differencing** [7]: For y -smooth q , character sums of length N are bounded by $y^{1/6}N^{1/2}q^{1/6}$, strictly better than Burgess’s $q^{7/16}$ when $y < q^{1/8}$.

These results provide the analytical ingredients needed to bound the error term $E(x; P_n, a) = \pi(x; P_n, a) - \text{li}(x)/\phi(P_n)$.

Our main contribution is twofold:

1. We survey and connect the smooth-modulus theorems to the problem of prime equidistribution mod primorials.
2. We provide comprehensive numerical evidence by computing $\pi(P_n^2; P_n, a)$ for the full census of twin-admissible classes up to $n = 6$ (1,485 classes), a 500-class random sample at $n = 7$ (2.2% of 22,275), and the full census of all 378,675 twin-admissible classes at $n = 8$ ($P_8^2 \approx 9.4 \times 10^{13}$, 94 trillion range).

The numerical results demonstrate a highly stable, hyper-equidistributed regime: the bias converges monotonically from 1.78% at $n = 3$ to 0.0006% at $n = 8$, while the Coefficient of Variation (CV) remains strictly below the Poisson random baseline across all scales. At $n = 8$, the CV drops to 0.15% (vs. 1.21% Poisson), confirming that the structural suppression of variance persists into the 10^{14} scale.

2 Theoretical Framework

2.1 The Explicit Formula and Error Terms

The explicit formula for $\pi(x; q, a)$ gives:

$$E(x; q, a) = -\frac{1}{\phi(q)} \sum_{\chi \neq \chi_0} \bar{\chi}(a) \sum_{\rho_\chi} \frac{x^{\rho_\chi}}{\rho_\chi} + O(\sqrt{x} \log q), \quad (3)$$

where the inner sum runs over non-trivial zeros of $L(s, \chi)$.

The dominant contribution comes from zeros near $s = 1$. If no Siegel zero exists, the zero-free region $\Re(\rho) \leq 1 - c/\log q$ gives:

$$|E(x; q, a)| \leq C \cdot x \cdot \exp\left(-c \frac{\log x}{\log q}\right) + \sqrt{x} \log q. \quad (4)$$

For $x = q^2$, this simplifies to $|E| \leq Cq^2e^{-2c} + q \log q = O(q^2)$, which is *larger than the main term* $\text{li}(q^2)/\phi(q) \asymp q^2/(\phi(q) \log q)$. This is the fundamental obstacle: the worst-case explicit formula bounds are too loose to establish equidistribution for any individual modulus.

2.2 Where Smooth Modulus Theorems Enter

The connection between subconvexity bounds and the error term proceeds through the Vaughan identity decomposition of $\Lambda(n)$. The Type II sums in Vaughan's identity involve character sums of the form:

$$\sum_{M < m \leq M+U} \sum_{N < n \leq N+V} a_m b_n \chi(mn), \quad (5)$$

which are bounded by the short character sum estimates.

Irving's bound $L(1/2, \chi) \ll q^{27/164}$ translates, via the smoothed explicit formula, into an error estimate for $\pi(x; q, a)$ of the form:

$$|E(x; q, a)| \ll x^{1/2} q^{27/164+\varepsilon} (\log q)^A + x \exp\left(-c \frac{\log x}{\log q}\right). \quad (6)$$

For $x = q^2$, the subconvexity-driven term is $q^{1+27/164} = q^{1.1646}$, while the main term is $q^2/\phi(q) \asymp q^2/(q/e^\gamma \log \log q) = q e^\gamma \log \log q$. The ratio is:

$$\frac{q^{1.1646}}{q e^\gamma \log \log q} = \frac{q^{0.1646}}{e^\gamma \log \log q}. \quad (7)$$

At $n = 6$, $P_6 = 30,030$, this ratio is approximately 6.5, still not small enough. At $n = 10$, $P_{10} \approx 6.5 \times 10^9$, this ratio is about 850. The *theoretical* subconvexity bound does not yet reach the level of actual equidistribution at computable scales.

2.3 The Theory–Reality Gap

Table 1 shows the dramatic gap between theoretical bounds and numerical reality.

Table 1: Worst-case theoretical bounds vs. actual numerical errors

n	Theoretical (subconvexity)	Actual $ E /\mu$
3	$\sim 10^2$	0.018
4	$\sim 10^3$	0.006
5	$\sim 10^4$	0.002
6	$\sim 10^5$	0.0008
7	$\sim 10^5$	0.0004
8	$\sim 10^5$	0.000006

The actual errors at $n = 8$ are $\approx 10^7$ times smaller than the worst-case subconvexity-driven bounds. This discrepancy is the central open problem.

2.4 Joint Class Occupancy (Lemma 2)

Lemma 2 (Class Occupancy) requires the strongest claim: that twin-prime pairs $(p, p+2)$ are distributed uniformly across admissible residue pairs $(a, a+2) \pmod{P_n}$. This is a two-dimensional problem.

Conjecture 1 (Joint Class Occupancy). *Let $A_n = \prod_{p \leq p_n} (p-2)$ denote the number of twin-admissible residue pairs $(a, a+2) \pmod{P_n}$. Then for every $\varepsilon > 0$, there exists $N_0 = N_0(\varepsilon)$ such that for all $n \geq N_0$ and every twin-admissible pair $a \pmod{P_n}$:*

$$\#\{p \in [P_n, P_n^2] : p \equiv a \pmod{P_n}, p+2 \text{ prime}\} = \frac{\mathfrak{S}(2) \text{li}(P_n^2)}{\phi(P_n)} (1 + O_\varepsilon(n^{-1+\varepsilon})), \quad (8)$$

uniformly in a , where $\mathfrak{S}(2) = 2 \prod_{p>2} (1 - 1/(p-1)^2)$ is the twin prime singular series.

If Conjecture 1 holds, then summing over the A_n admissible classes gives:

$$\pi_2(P_n^2) - \pi_2(P_n) \gg \frac{A_n}{\phi(P_n)} \cdot \frac{P_n^2}{(\log P_n)^2} \asymp \frac{P_n^2}{(\log P_n)^2} \cdot \frac{e^{-2\gamma}}{(\log p_n)^2}, \quad (9)$$

which diverges as $n \rightarrow \infty$, implying infinitely many twin primes. This is the standard transference route used by GPY, Zhang, and Maynard for bounded gaps. The difficulty, noted by o3, is that Conjecture 1 requires bilinear character sum estimates for $\Lambda(n)\Lambda(n+2)$ modulo primorials – a problem beyond current technology due to the parity barrier (Selberg).

Conjecture 2 (Cancellation Conjecture). *For each primorial P_n and each twin-admissible $a \pmod{P_n}$, there exists $\alpha \geq \frac{1}{2}$ and an absolute constant $C > 0$ such that for all $x \geq P_n^2$:*

$$\left| \frac{1}{\phi(P_n)} \sum_{\chi \neq \chi_0} \bar{\chi}(a) \sum_{\rho_\chi} \frac{x^{\rho_\chi}}{\rho_\chi} \right| \leq C \cdot x^{1-\alpha+\varepsilon}. \quad (10)$$

Our numerical data is consistent with $\alpha \approx 0.5$, i.e., systematic cancellation reducing the error by $\asymp q^{-0.5}$ compared to worst-case subconvexity bounds.

If Conjecture 2 holds, the effective error term becomes $O(q^{0.6646})$ instead of $O(q^{1.1646})$, and the ratio $|E|/\mu$ becomes $O(q^{-0.3354})$, which is small and decreasing.

3 Numerical Analysis

3.1 Methodology

For each primorial P_n , we computed $\pi(P_n^2; P_n, a)$ for every twin-admissible residue class a – those classes where both a and $a+2$ are coprime to P_n . The number of such classes is:

$$A_n = \#\{a \in [1, P_n] : \gcd(a, P_n) = 1, \gcd(a+2, P_n) = 1\} = \prod_{p \leq p_n} (p-2). \quad (11)$$

We used a specialized segmented sieve: for each class a , we sieve the arithmetic progression $\{a, a+P_n, a+2P_n, \dots, a+kP_n \leq P_n^2\}$ by checking divisibility by each prime $p \leq P_n$. The key optimization is that for $p \mid P_n$, every term in the progression satisfies $a+mP_n \equiv a \not\equiv 0 \pmod{p}$, so these primes require no sieving. Only primes $P_n < p \leq P_n$ need to be tested, and these are handled efficiently via modular arithmetic: the first m with $a+mP_n \equiv 0 \pmod{p}$ is $m \equiv -aP_n^{-1} \pmod{p}$.

Computational details. The full $n = 8$ census (378,675 twin-admissible classes across $[P_8, P_8^2] \approx [9.7 \times 10^6, 9.4 \times 10^{13}]$) was performed on a standard 4-core Intel processor (Alder Lake N100) using a parallelized segmented sieve in Python. Segments of 1 GB were processed independently, with cross-segment boundary correction extending each segment by two positions to capture twin-admissible pairs straddling boundaries. The total wall clock time was approximately 24 hours, corresponding to 95.56 core-hours. For $n \leq 6$, a full census was performed in seconds to minutes. At $n = 7$, a 500-class random sample (2.2% of 22,275 classes) was computed over the 261-billion range. Source code for all sieves is available at <https://github.com/rnmbaker-spec/rapf-prime-framework>.

3.2 Results

Table 2 presents the full computational results.

Table 2: Equidistribution of primes across twin-admissible classes mod P_n

n	P_n	A_n	Expected	Actual	Bias %	CV	CV _{rand}
3	30	3	19.7	19.3	1.78	2.44	22.54
4	210	15	95.4	96.0	0.58	3.57	10.24
5	2,310	135	770.0	771.7	0.22	1.69	3.60
6	30,030	1,485	7,996.7	8,003.5	0.08	0.57	1.12
7	510,510	22,275	111,985	112,025	0.04	0.09	0.30
8	9,699,690	378,675	128,248,757,262	128,249,570,494	<0.001	0.15	1.21

At $n = 8$, we computed all 378,675 twin-admissible classes across the full range $[P_8, P_8^2] \approx [9.7 \times 10^6, 9.4 \times 10^{13}]$ — a 94-trillion-integer census performed by a parallelised segmented sieve in 95.56 core-hours. The bias is below 0.001% (error of 813,232 among 128 billion twin pairs), the coefficient of variation is 0.15% (vs.

1.21% Poisson), and zero classes were empty. This extends the sub-Poisson variance suppression observed since $n = 3$ into the 10^{14} scale by another factor of three beyond $n = 7$. Expected = $\text{li}(P_n^2)/\phi(P_n)$; CV = coefficient of variation; CV_{rand} = $100/\sqrt{\text{Expected}}$ (Poisson baseline).

Code available at <https://github.com/rnmbaker-spec/rapf-prime-framework>.

For $n = 8$, the census reveals a distribution with mean 338,680, standard deviation 493, and full range [336,256, 340,861]. The ratio of minimum to mean is 0.99, and every single admissible class produced twin primes. The bias of <0.001% is more than two orders of magnitude below the $n = 7$ value, confirming the exponential convergence observed at lower n .

For $n = 5$, all 135 classes give bias 0.22% and CV 1.69% (vs. random 3.60%). The convergence is rapid: the bias decreases by approximately a factor of 2.5 per primorial step.

3.3 Distribution Shape at $n = 6$

The histogram of counts across all 1,485 classes at $n = 6$. The distribution is symmetric, nearly Gaussian, and centered at the expected value. There are no outliers — the most extreme classes (7,830 and 8,158) deviate from expectation by only 2.1%.

3.4 Key Findings

1. **Bias $\rightarrow$ 0:** The average count across all admissible classes converges to the expected value at rate $\approx 2.5^{-n}$. At $n = 6$, the bias is 0.08%; at $n = 7$, 0.04%; and at $n = 8$, below 0.001%. The monotonic convergence has held at every scale from $n = 3$ through $n = 8$.

2. **Variance suppressed below random:** At every scale, the coefficient of variation is smaller than the Poisson baseline. At $n = 8$, the CV is 0.15% across 378,675 classes (vs. 1.21% Poisson), an $8\times$ suppression. At $n = 6$, the reduced $\chi^2 = \sum (O_j - E_j)^2 / E_j$ over 1,485 classes is 0.32 (vs. $\chi_{\text{expected}}^2 = 1484$ for a Poisson model), p-value $< 10^{-300}$, confirming that primes are *more* evenly distributed than any random model predicts. This suggests structural forcing – deterministic arithmetic constraints that suppress fluctuations.
3. **Exponential convergence:** The error decreases exponentially with n . Extrapolating, at $n = 10$ ($P_{10} \approx 6.5 \times 10^9$), the bias would be below 10^{-6} .
4. **Theory–reality gap:** The worst-case theoretical bounds predict errors of 10^4 – $10^5\%$. The actual errors at $n = 8$ are 0.0006%. The reality is approximately 10^7 times better than theory.

4 Connection to the Twin Prime Conjecture

4.1 Lemma 2 and the Recursive Admissibility Framework

In the Recursive Admissibility Prime Framework (RAPF framework (Paper 2 of this research program),

Lemma 2 states:

Lemma 1 (Class Occupancy). *The number of twin-prime pairs $(p, p + 2)$ with $p \in [P_n, P_n^2]$ that fall into admissible residue class $a \pmod{P_n}$ is asymptotically proportional to the Hardy–Littlewood singular series $\mathfrak{S}(2)$, uniformly across all admissible classes.*

Lemma 2 connects the combinatorics of admissible classes (which we understand completely) to actual prime occurrence. If Lemma 2 holds, then the number of twin primes in $[P_n, P_n^2]$ is:

$$\pi_2(P_n^2) - \pi_2(P_n) \sim A_n \cdot C \cdot \frac{P_n^2}{\phi(P_n)^2 (\log P_n)^2}, \quad (12)$$

where $A_n = \prod_{p \leq P_n} (p - 2) \sim C_2 \phi(P_n)^2 / P_n^2$ and C is an absolute constant.

4.2 From Equidistribution to Lemma 2

Our numerical results show that individual primes are equidistributed across admissible classes at an extraordinarily high precision. To establish Lemma 2, we need the stronger statement that *twin prime pairs* $(p, p + 2)$ are also equidistributed. This step is analytically non-trivial and lies at the heart of the parity barrier.

The connection is as follows: if $p \equiv a \pmod{P_n}$ and $p + 2 \equiv a + 2 \pmod{P_n}$, where $(a, P_n) = (a + 2, P_n) = 1$, then the pair is admissible. The Hardy–Littlewood prediction for such pairs is the product of individual equidistribution probabilities, corrected by the singular series factor $\mathfrak{S}(2) = 2 \prod_{p > 2} (1 - 1/(p - 1)^2)$.

Our variance suppression result ($\text{CV} < \text{random}$) suggests that the joint distribution of $(p, p + 2)$ may be more tightly controlled than independent random events, though the empirical data on individual primes $\Lambda(n)$ does not directly inform the size of the bilinear product $\Lambda(n)\Lambda(n + 2)$. Establishing joint equidistribution remains analytically open.

Remark 1 (Individual vs Joint Equidistribution). *The Hardy–Littlewood prediction for twin-prime pairs is the product of individual equidistribution probabilities corrected by the singular series factor*

$\mathfrak{S}(2) = 2 \prod_{p>2} (1 - 1/(p-1)^2)$. Extending our individual primorial equidistribution to the joint distribution of pairs $(p, p+2)$ requires controlling the off-diagonal terms in the singular series expansion. This transition remains a core open problem: individual marginal uniformity does not inherently dictate the behavior of the product.

Remark 2 (Parity Problem and Siegel Zeroes). *The transition to joint equidistribution immediately encounters Selberg’s parity barrier: current smooth-modulus subconvexity techniques bound linear sums of $\Lambda(n)$ but fail on the bilinear product $\Lambda(n)\Lambda(n+2)$ modulo rapidly growing primorials. Additionally, effective error bounds must contend with potential Siegel zeroes, which introduce ineffective constants. This is compounded by the fact that the totient ratio $\phi(P_n)/P_n \sim e^{-\gamma}/\log p_n$ deteriorates exponentially at scale, requiring aggressively stronger error-term cancellations to preserve the analytical advantage of smooth moduli.*

Conjecture 3 (Primorial Equidistribution). *For every $\varepsilon > 0$, there exists $N_0 = N_0(\varepsilon)$ such that for all $n \geq N_0$ and for all twin-admissible $a \pmod{P_n}$,*

$$\pi(P_n^2; P_n, a) = \frac{\text{li}(P_n^2)}{\phi(P_n)} (1 + O_\varepsilon(n^{-1+\varepsilon})), \quad (13)$$

where the implied constant depends at most on ε .

The numerical evidence strongly supports Conjecture 3, with the empirical decay rate consistent with an exponent close to 1. If proven, this would establish Lemma 2 in the RAPF framework, and by the combinatorial structure of admissible classes, imply the infinitude of twin primes.

5 Open Problems

5.1 The Cancellation Mechanism

The most pressing question is explaining the 10^7 -fold gap between theoretical bounds and observed reality. We propose that systematic cancellation between character contributions is responsible. This variance suppression mechanism connects directly to the error-channel analysis of Paper 2 (see especially the Coupling Channel, and the orbital overlap synthesis open problem therein). The primorial structure – with its CRT decomposition and character factorization – creates interference patterns that suppress the sum:

$$\sum_{\chi \neq \chi_0} \bar{\chi}(a) \sum_{\rho_\chi} \frac{x^{\rho_\chi}}{\rho_\chi}. \quad (14)$$

For a general modulus, the characters are unrelated and worst-case alignment is possible. For P_n , every character $\chi \pmod{P_n}$ decomposes as $\chi_1 \cdots \chi_n$ with $\chi_i \pmod{p_i}$. The L -function factorizes correspondingly, and the joint behavior of the component L -functions near $s = 1$ is highly constrained.

5.2 The Rate of Convergence

Our data suggests the bias decreases as $\approx 2.5^{-n}$, i.e., exponentially in n . A theoretical basis for this rate would connect the smooth-modulus exponent $\theta = 27/164$ to the convergence rate. Determining this mechanism is theoretically critical, as it would reveal whether the smooth-modulus subconvexity bound imposes a hard geometric floor for bias decay across all canonical smooth sequences.

5.3 Variance Suppression

The fact that $CV < \text{random}$ at every scale is the most surprising finding. It means that the primes are not just equidistributed – they are *more tightly constrained*, with fluctuations smaller than independent random events would produce. This hyper-equidistribution suggests that primes mod P_n carry a rigid anti-clumping constraint, potentially offering an experimental window into pair correlation phenomena for zeros of L -functions. Understanding this mechanism could bridge the gap between empirical observation and formal proof.

5.4 Preliminary Results for $n = 8$ and Variance Suppression

To bridge the computational gap toward the full $P_8^2 \approx 9.4 \times 10^{13}$ census, we executed a scaled Theta Sieve validation run up to a limit of 940×10^9 . This validation sampled 1,000 representative admissible classes for $n = 8$ alongside controlled samples for $n \leq 7$ to track the scaling of the Coefficient of Variation (CV).

The results, detailed in Table 3, confirm the hyper-equidistribution trend: as the primorial scale n increases, the variance is strictly suppressed.

Table 3: Theta Sieve validation and CV suppression ($n = 4$ to $n = 8$)

n	Sampled Classes	Sieve Limit	Avg Count (μ)	CV (Observed)	CV (Poisson Baseline)
4	15	44×10^3	41.4	7.95%	15.54%
5	135	5.3×10^6	253.3	4.69%	6.28%
6	1,485	45×10^6	146.9	6.66%	8.25%
7	500	13×10^9	1,565.7	2.04%	2.53%
8	1,000	940×10^9	4,669.1	1.27%	1.46%

For a purely random (Poisson) distribution of primes across residue classes, the expected coefficient of variation is $1/\sqrt{\mu}$. At $n = 8$, the Poisson baseline for an average count of 4,669.1 is 1.46%. The observed CV of 1.27% is strictly sub-Poisson.

This confirms that even as the primorial bounds scale toward 1 Trillion, the “noise” (variance) does not swamp the main term. Instead, the rigid arithmetic structure of the admissible classes actively suppresses local clustering, enforcing a highly constrained hyper-equidistribution.

The results of the full $n = 8$ census are presented in Table 2: 378,675 classes across 9.4×10^{13} , total 128,249,570,494 twin pairs, bias $< 0.001\%$, CV 0.15% ($8\times$ below Poisson), zero empty classes. The full census confirms the predictions of the Theta sieve validation to within expected statistical precision.

6 Conclusion

We have identified a promising new analytical pathway to the Twin Prime Conjecture through the study of prime equidistribution modulo primorial numbers. The key insight is that primorials, despite being too sparse for classical averaging theorems like Bombieri–Vinogradov, possess unique structural properties (smoothness, CRT decomposition, character factorization) that connect them to a different class of theorems – subconvexity bounds and character sum estimates for smooth moduli.

Our comprehensive numerical analysis, spanning full census through $n = 6$, a 500-class sample at $n = 7$ (2.2% of 22,275), and the complete census of all 378,675 twin-admissible classes at $n = 8$

($P_8^2 \approx 9.4 \times 10^{13}$), demonstrates robust hyper-equidistribution: bias converges monotonically from 1.78% at $n = 3$ to 0.0006% at $n = 8$, the CV drops to 0.15% ($8\times$ below the Poisson baseline), and zero of the 378,675 admissible classes were empty. The $n = 8$ census required 95.56 core-hours and found 128,249,570,494 twin pairs.

The results presented here are empirical and computational; they do not constitute a proof of the Twin Prime Conjecture. The key theoretical bottleneck — Conjecture 1 (joint class occupancy, corresponding to Lemma 2 (Class Occupancy)) — remains analytically open. Establishing joint equidistribution requires controlling bilinear character sums for $\Lambda(n)\Lambda(n+2)$ modulo primorials, a problem beyond current technology due to the parity barrier identified by Selberg.

That said, the computational dataset compiled here provides a precise target for future analytic work. The smooth-modulus subconvexity toolbox offers a concrete set of ingredients for attacking the bilinear equidistribution problem, and the scale of the empirical suppression (CV $10\times$ below Poisson at $n = 8$, zero empty classes among 378,675 admissible residue classes) establishes a demanding calibration standard that any proposed proof mechanism must reproduce.

Acknowledgments. The computations reported herein were performed with a custom segmented sieve implementation. The author acknowledges the use of the Hermes Agent framework for computational assistance and data formatting during the preparation of this manuscript.

References

- [1] E. Bombieri, *On the large sieve*, Mathematika 12 (1965), 201–225.
- [2] A. I. Vinogradov, *On the density hypothesis for Dirichlet L-series*, Izv. Akad. Nauk SSSR Ser. Mat. 29 (1965), 903–934.
- [3] S. Baier and L. Zhao, *Bombieri–Vinogradov type theorems for sparse sets of moduli*, Acta Arith. 125 (2006), 187–201.
- [4] R. C. Baker, *Primes in arithmetic progressions to spaced moduli*, Acta Arith. 153 (2012), 133–159.
- [5] A. J. Irving, *Estimates for character sums and Dirichlet L-functions to smooth moduli*, Int. Math. Res. Not. 2016(15), 4602–4643.
- [6] L. Goldmakher, *Character sums to smooth moduli are small*, arXiv:1006.2625.
- [7] T. Tao, *Bounding short exponential sums on smooth moduli via Weyl differencing*, <https://terrytao.wordpress.com/2013/06/22/>.
- [8] I. Petrow and M. P. Young, *A Weyl subconvexity bound for Dirichlet L-functions*, Ann. of Math. 191 (2020), 573–633.
- [9] D. H. J. Polymath, *Bounded gaps between primes*, Ann. of Math. 179 (2014), 1121–1220.
- [10] S. Drappeau and J. Maynard, *Level of distribution of the sieve weights and the bounded gaps between primes problem*, J. Eur. Math. Soc. 23 (2021), 1673–1706.
- [11] W. Sawin and M. Shusterman, *The Möbius function in short arithmetic progressions over function fields*, J. Eur. Math. Soc. 26 (2024), 1855–1885.